



LET'S PLAY!

W

O

R

D

L

E

A DAILY WORD GAME

A Reinforcement learning Approach

Poornima Devi Krishnasamy Karthikeyan



HOW TO PLAY

Guess the Wordle in 6 tries.

1 Each guess must be a valid 5-letter word.

2 Color of Titles will change to show how close your guess was to the word

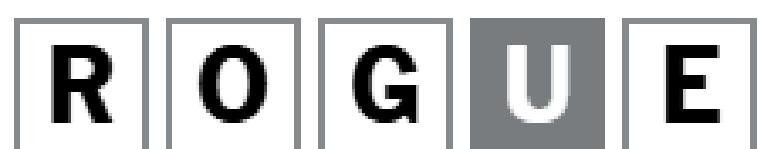
Examples



W is in the word and in the correct spot.

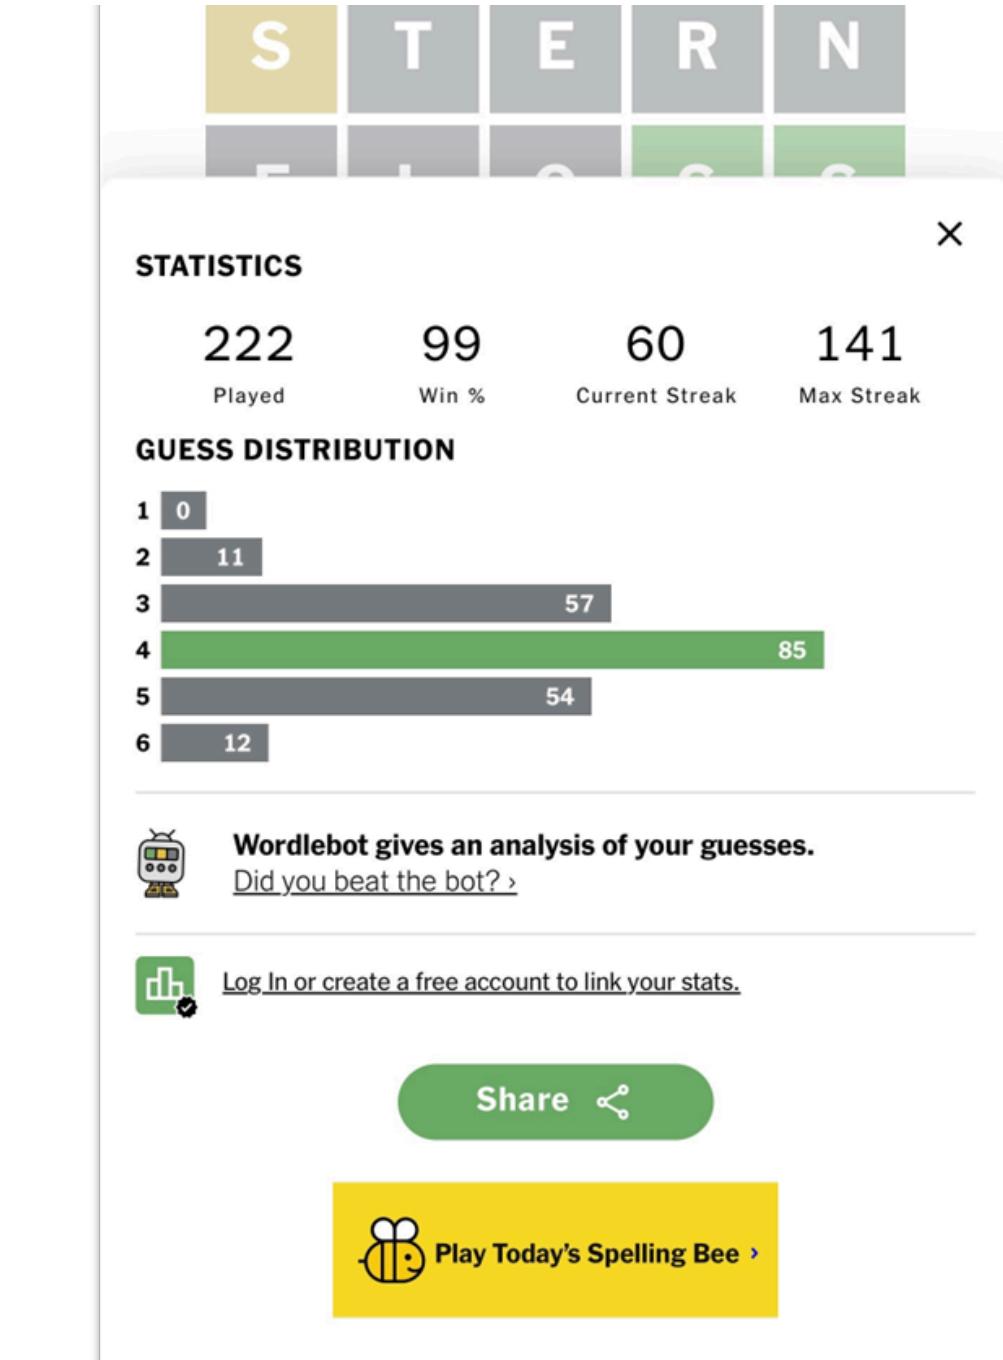
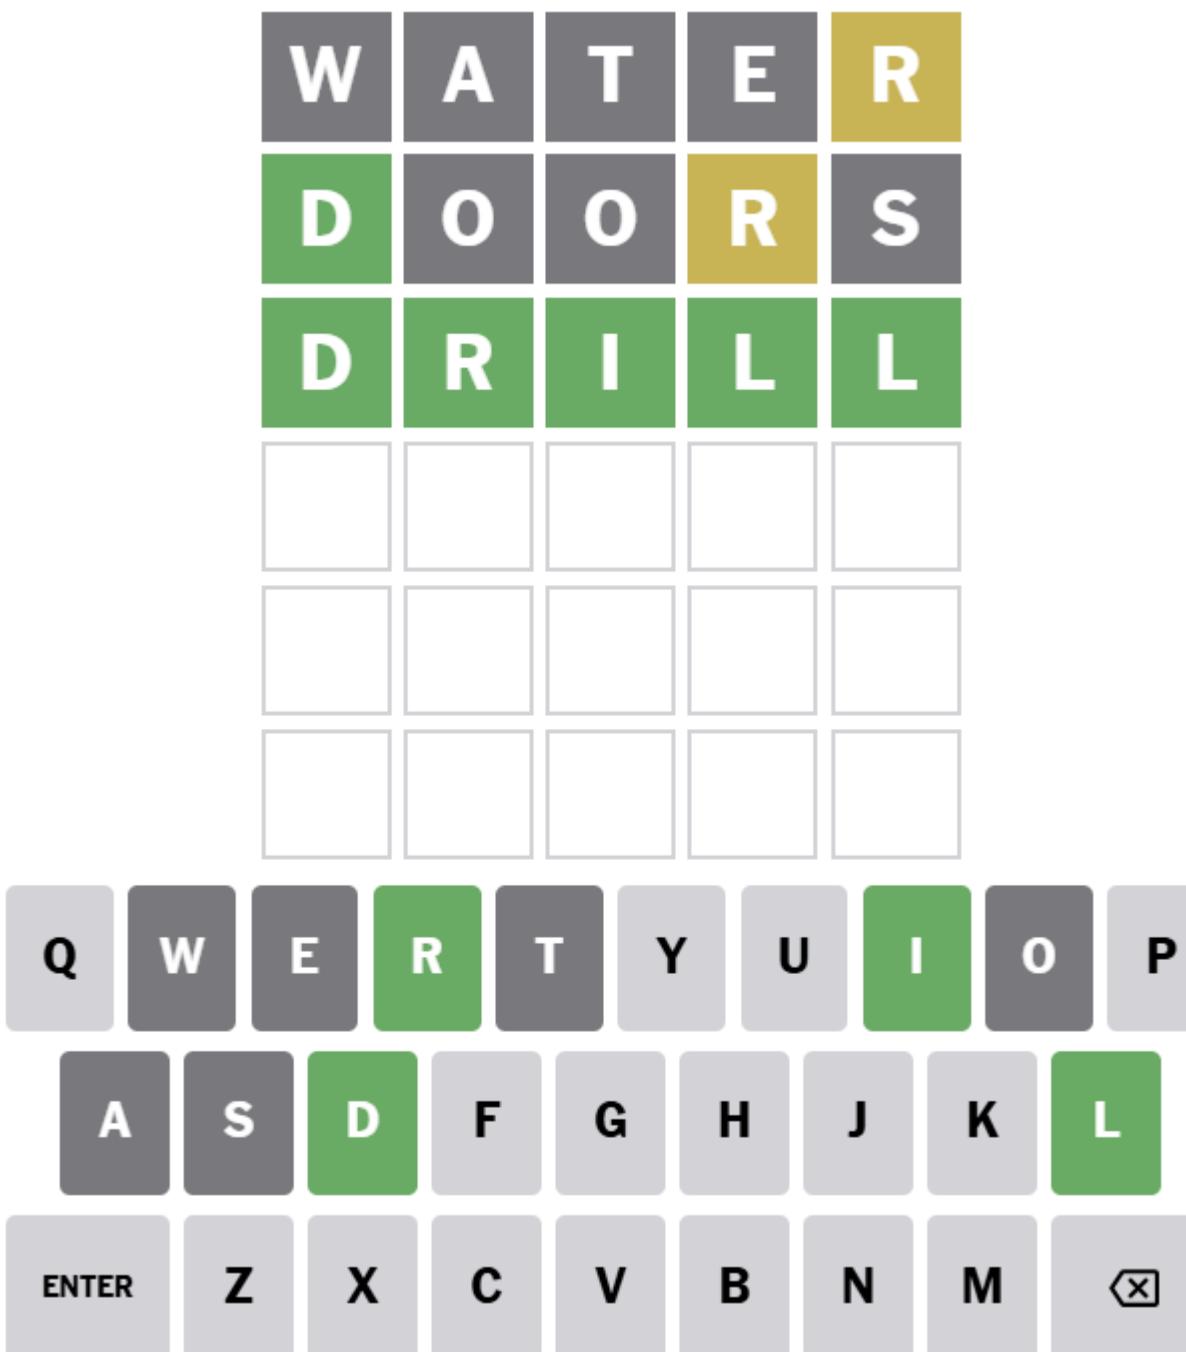


I is in the word but in the wrong spot.



U is not in the word in any spot.

GAME PLAY





CHALLENGE

GOAL → Guess 5-letter word in 6 attempts

Receive color-coded feedback:

- Green: Correct letter, correct position
- Yellow: Correct letter, wrong position
- Gray: Letter not in word

OBJECTIVE → Maximize win rate • Minimize average guesses



VOCABULARY LIST

ACTUAL TEXT LIST

12972 WORDS

- 1 cigar
- 2 rebut
- 3 sissy

.....

- 12969 zygal
- 12970 zygon
- 12971 zymes
- 12972 zymic

SOLUTION TEXT LIST

2315 WORDS

- 1 aback
- 2 abase
- 3 abate

.....

- 2312 youth
- 2313 zebra
- 2314 zesty
- 2315 zonal

→ TARGET



MOTIVATION TO USE RL

- Sequential problem: each guess affects future states.
- Built-in reward signal (colour feedback).
- Adapts easily to rule changes (word length, attempts).
- Avoids brute-force search , fixed rule based system or trying to define what the best guess is through entropy, information gain, likelihood, etc.



EVVIRONMENT

- Game state representation
- Feedback mechanism
- Reward function
- Implementation



GAME STATE REPRESENTATION

State representation: 725-dimensional vector capturing the full game state, including:

- Encoded feedback from previous guesses
- Letter constraints (present/absent/position)
- Remaining valid word structure

Action representation: Selecting a word from the dictionary, with masking to prevent repeated guesses.



FEEDBACK MECHANISM

After each guess, the environment returns a 5-element feedback vector:

- 2 → Green (correct letter, correct position)
- 1 → Yellow (correct letter, wrong position)
- 0 → Gray (letter not in the target word)

This is the state information. Agent uses this to update the state representation.



REWARD FUNCTION

1. Solve the word (Correct guess):

- Reward = 10 points
- Bonus = 2 points for each remaining turn

Encourages the agent to solve the word as quickly as possible.

2. Informative letters:

- Green letter (correct position) = +0.8
- Yellow letter (correct letter, wrong position)= +0.3

Rewards guesses that give useful information.

3. Unhelpful guesses:

- All gray letters (no info) = -0.5

Discourages wasting guesses on uninformative words.

4. Fail to solve (turns exhausted):

- Reward = -5

Penalizes failure to solve the word.



FEEDBACK MECHANISM

After each guess, the environment returns a 5-element feedback vector:

- 2 → Green (correct letter, correct position)
- 1 → Yellow (correct letter, wrong position)
- 0 → Gray (letter not in the target word)

This is the state information. Agent uses this to update the state representation.

EXPLORATION

Adaptive Temperature:

- Starts high (~1.5) → explore many words, discover strategies
- Gradually decreases to low (~0.3) → exploit best known words

```
def masked_softmax(logits, mask, episode, total_episodes):
    adaptive_temp = max(0.3, 1.5 - (episode / total_episodes) * 1.2)
    logits = logits / max(adaptive_temp, 1e-6)
    logits = logits.masked_fill(mask == 0, -1e9)
    return F.softmax(logits, dim=-1), adaptive_temp
```

Behavior:

- High Temp: Tries diverse words, avoids local optima
- Low Temp: Focuses on high-probability words, maximizes success



ALGORITHM CHOSEN

Environment formulation reveals that state space and action space are large. Not feasible to apply tabular methods. Hence, the following algorithm was implemented:

- **Advantage Actor Critic** → A2C handles large state and action spaces efficiently by learning both the policy and value, enabling smart and stable Wordle strategies.

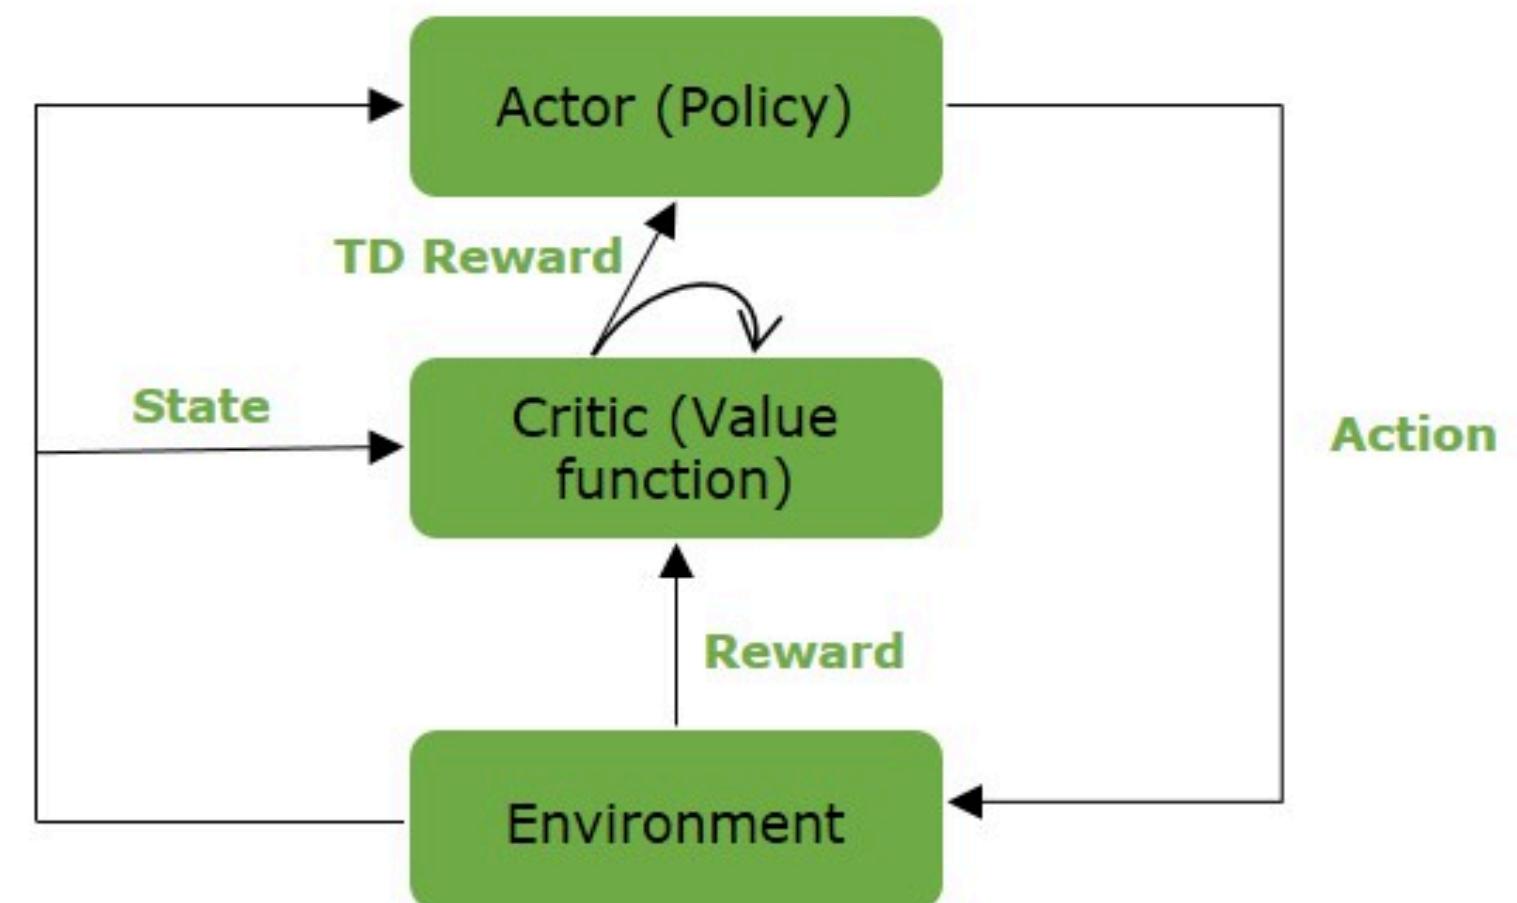
ADVANTAGE ACTOR CRITIC

Actor → Selects actions based on the policy to **maximize rewards**, continually refining it to adapt to the environment.

Critic → Evaluates actor's actions, providing feedback to guide towards higher returns and improve learning.

The **advantage function**, $A(s,a)$, measures the advantage of taking action a in state s over the expected value of the state under the current policy.

$$A(s,a) = Q(s,a) - V(s)$$



RESULTS

When training the model for 5000 epochs, it took approximately 1 hour.

Trained Agent

- 93% success rate
- 4.30 average guesses per win

Target: SHIED

Turn 1: CHITS

Turn 2: PHISH

Turn 3: SHIER

Turn 4: SHIED

SOLVED in 4 turns!

RESULTS : BEST STARTING GUESS

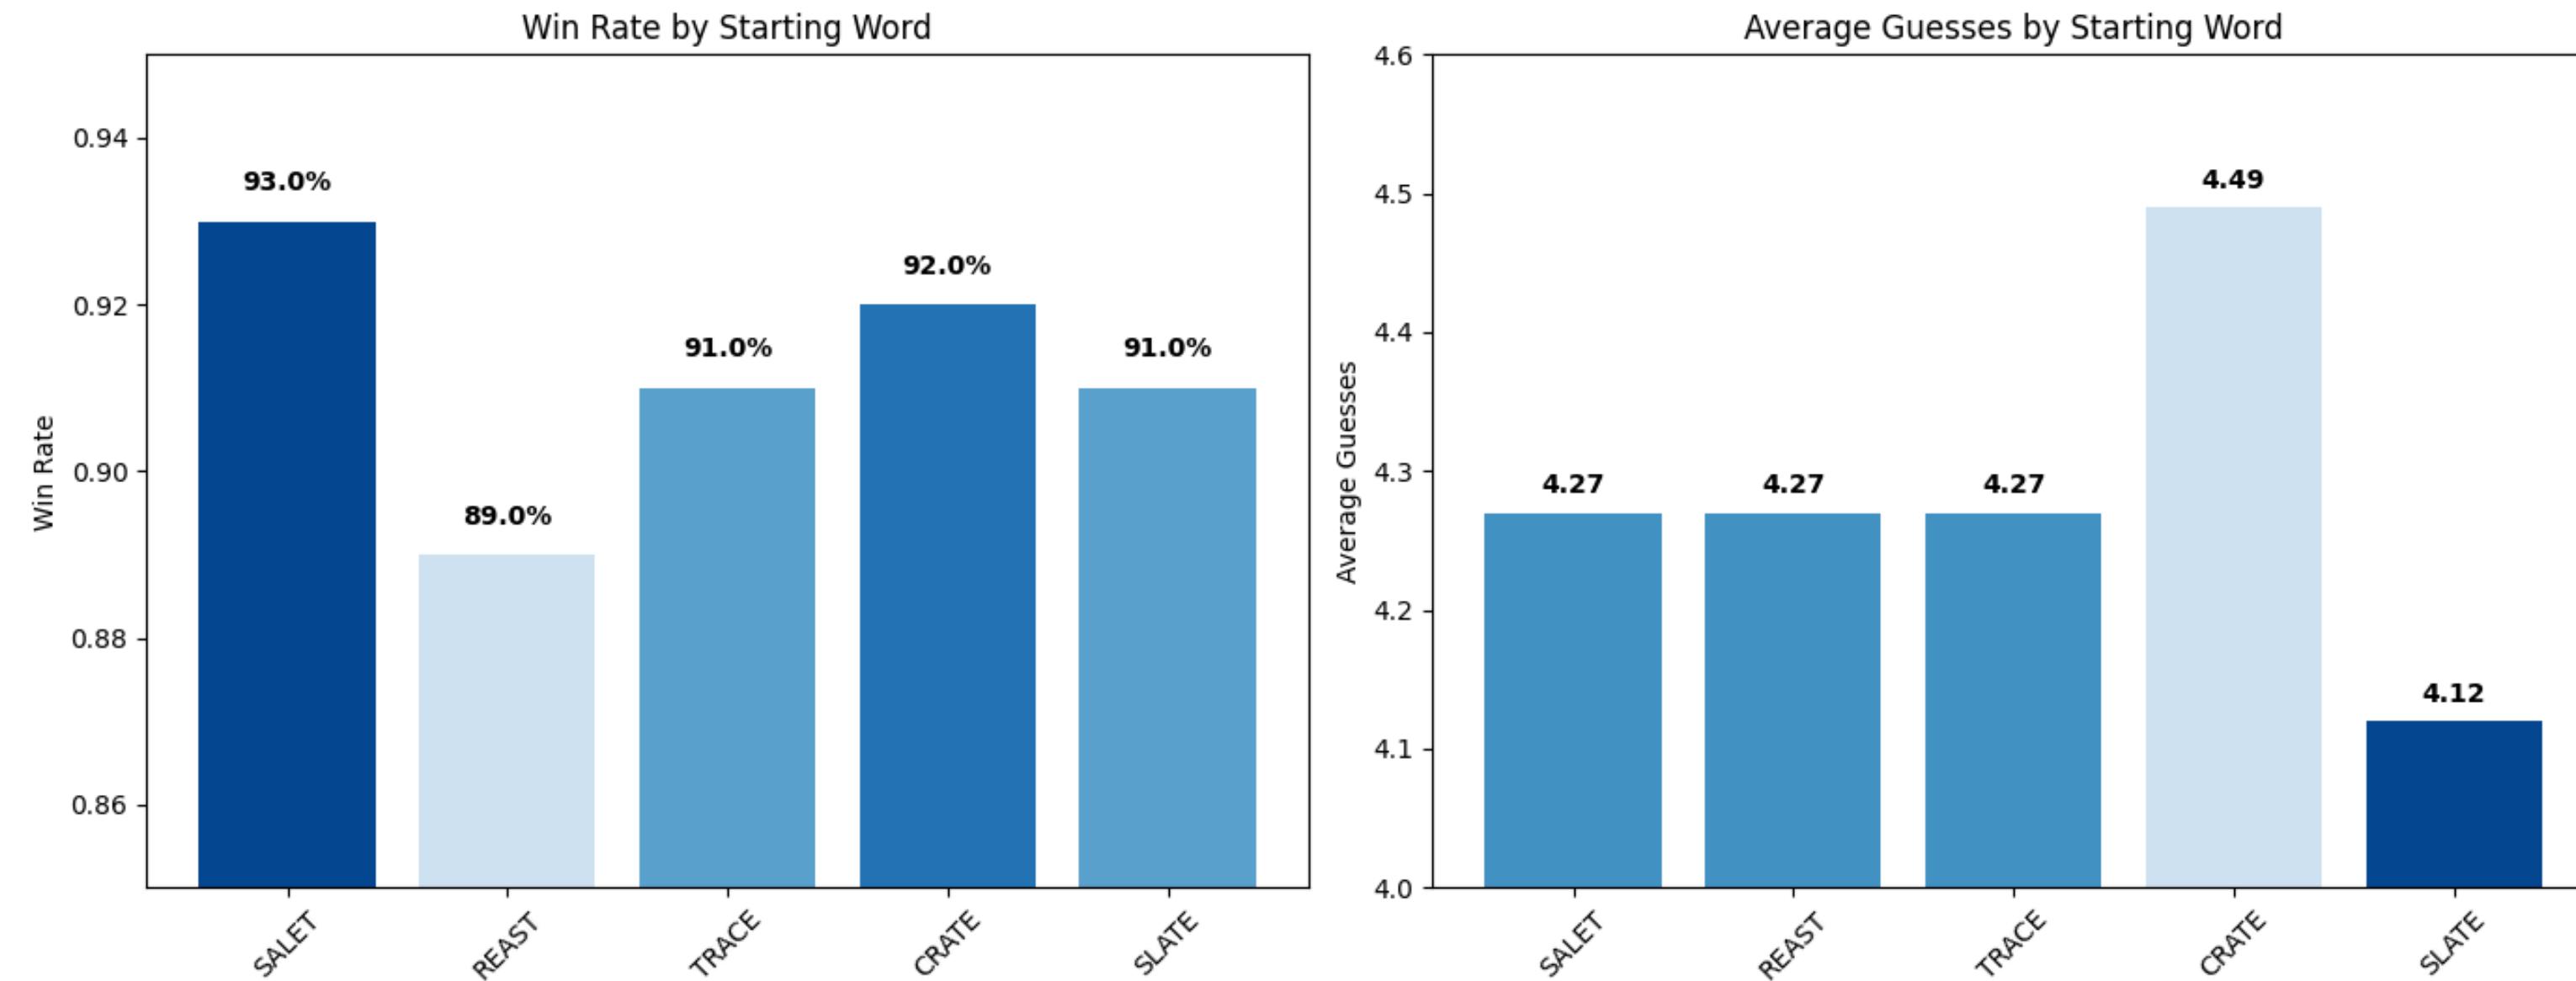
HOW DO YOU PICK A WORD THAT WILL GET TO THE SOLUTION WITH IN 6 GUESSES?

Initial Guess	Expected Number of Guesses
SALET	3.42117
REAST	3.42246
TRACE	3.42376
CRATE	3.42376
SLATE	3.42462

Table 1 : listing the five initial guesses in wordle yielding the lowest expected number of guesses to win. If an optimal policy is followed

Reference : https://auction-upload-files.s3.amazonaws.com/Wordle_Paper_Final.pdf

RESULTS : AGENT WITH BEST STARTING GUESS



SALET is the best starting word with 93 % win rate and 4.12 average guess



RESULTS

When training the model for 5000 epochs, it took approximately 1 hour.

Trained Agent → CHITS

- 93% success rate
- 4.30 average guesses per win

Target: SHIED

Turn 1: CHITS

Turn 2: PHISH

Turn 3: SHIER

Turn 4: SHIED

SOLVED in 4 turns!

Agent with best starter → SALET

- 93% success rate (100 games)
- 4.12 average guesses per win

Target: STOOD

Turn 1: SALET (FORCED STARTER)

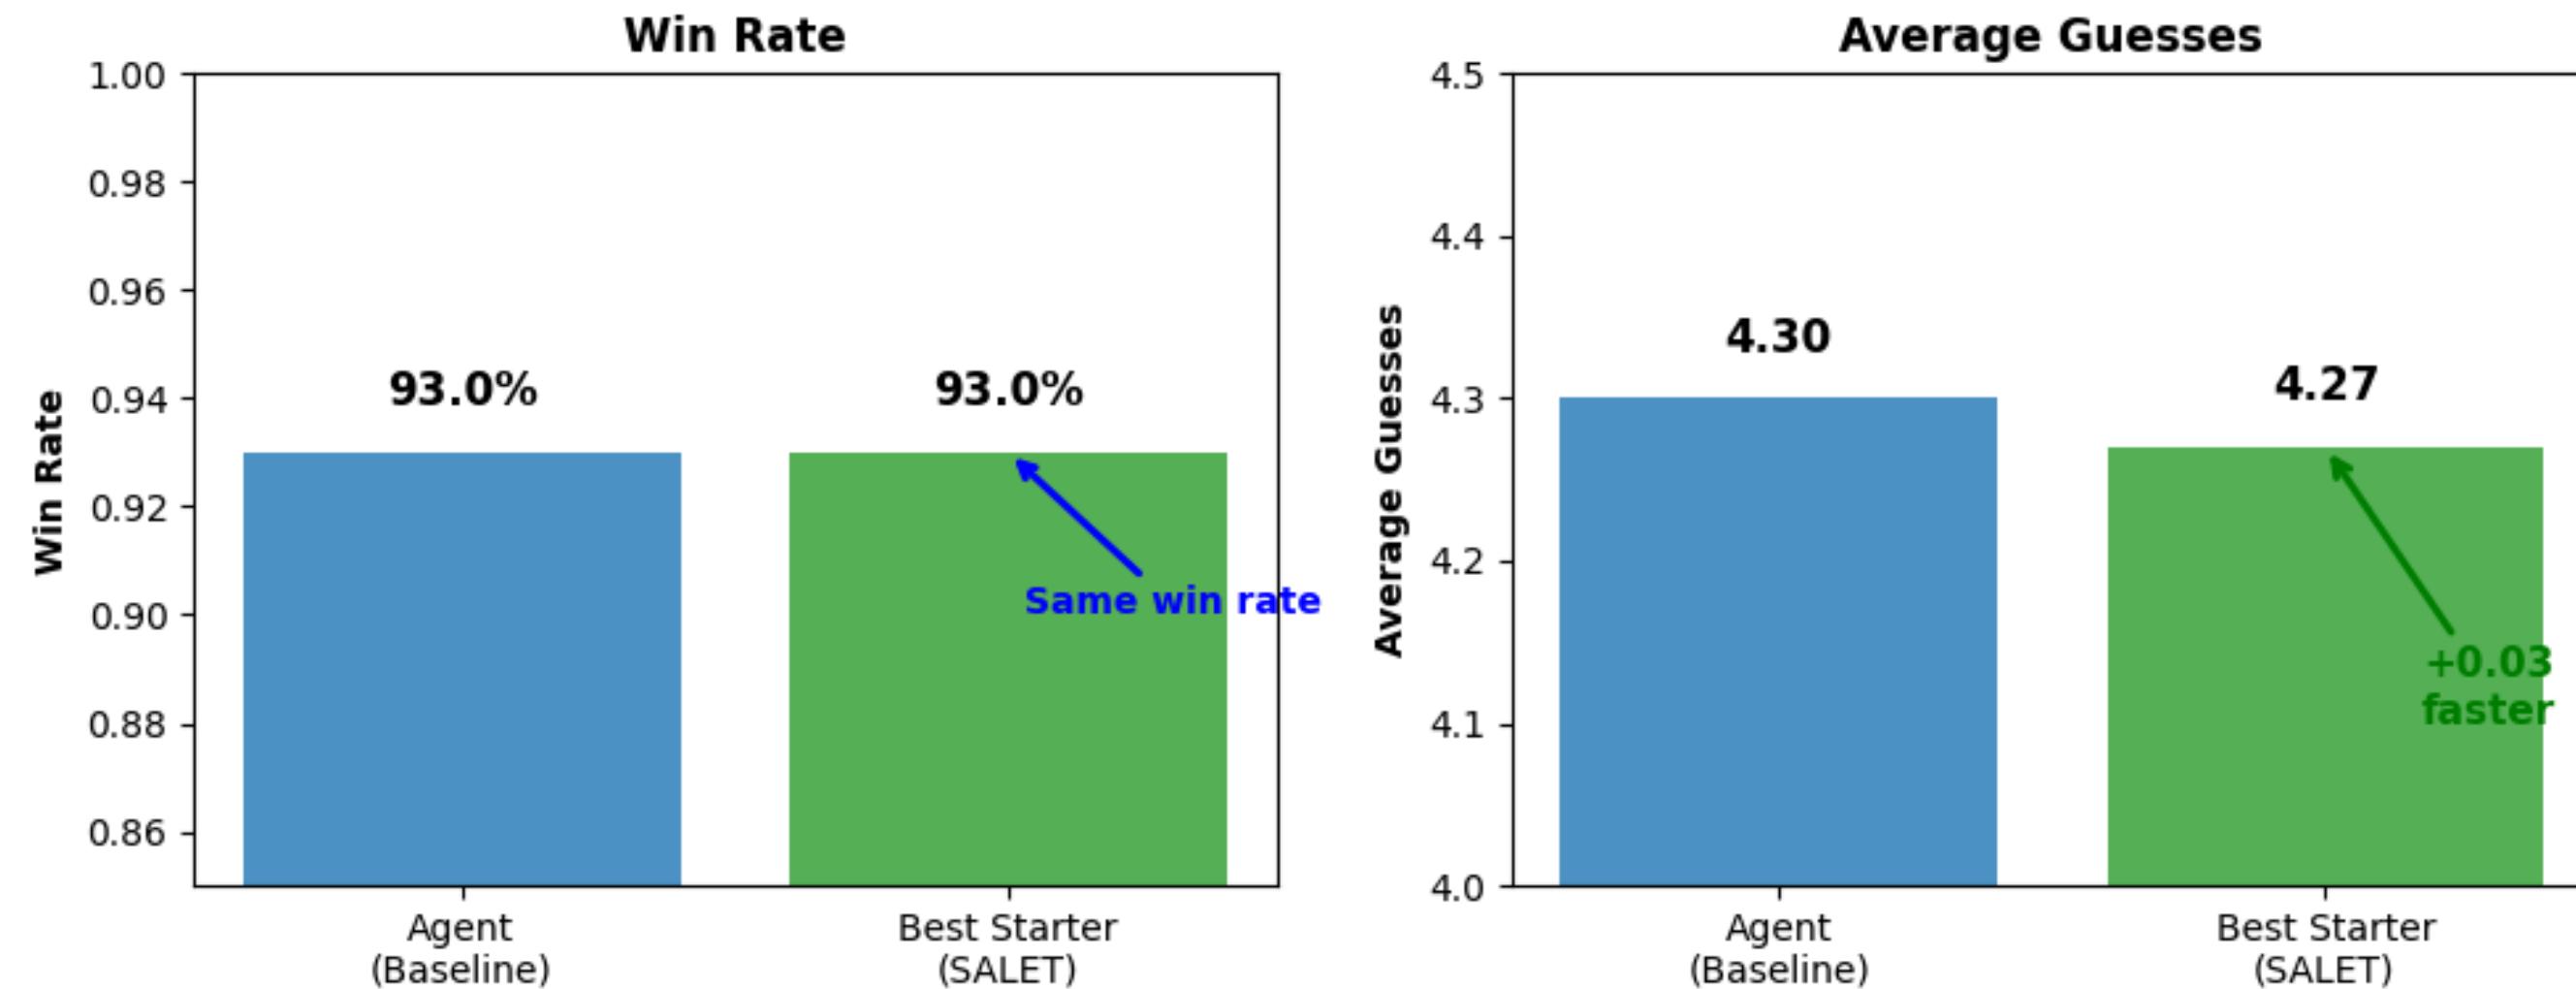
Turn 2: STROY

Turn 3: STOOK

Turn 4: STOOD

SOLVED in 4 turns!

RESULTS



comparison of Baseline agent and agent with best starting word



THANK YOU